

# Lab Programs:

## Week 8: Heart Disease Program

Aim: To analyze a heart disease dataset and visualize relevant insights.

Programs:

### 1. Program 1: Loading Heart Disease Dataset

- Procedure: Load the heart disease dataset and display its structure.

- Code:

```
import pandas as pd

df = pd.read_csv('heart_disease.csv')

print(df.head())
```

- Expected Result: A preview of the heart disease dataset.

### 2. Program 2: Summary Statistics

- Procedure: Generate summary statistics of the dataset.

- Code:

```
print(df.describe())
```

- Expected Result: Summary statistics including count, mean, std, min, and max values.

### 3. Program 3: Correlation Matrix

- Procedure: Create a correlation matrix and visualize it using a heatmap.

- Code:

```
import seaborn as sns

import matplotlib.pyplot as plt

correlation = df.corr()

sns.heatmap(correlation, annot=True, cmap='coolwarm')

plt.title("Correlation Matrix for Heart Disease Dataset")

plt.show()
```

- Expected Result: A heatmap displaying the correlation matrix for the features in the dataset.

#### 4. Program 4: Visualizing Age Distribution

- Procedure: Visualize the age distribution of patients with heart disease.

- Code:

```
sns.histplot(df['age'], bins=20, kde=True)

plt.title("Age Distribution of Heart Disease Patients")

plt.xlabel("Age")

plt.ylabel("Frequency")

plt.show()
```

- Expected Result: A histogram with a KDE curve showing the age distribution.

#### 5. Program 5: Classification of Heart Disease

- Procedure: Train a simple logistic regression model to classify heart disease.

- Code:

```
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score

X = df.drop('target', axis=1)

y = df['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LogisticRegression(max_iter=1000)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
```

- Expected Result: An accuracy score indicating the model's performance.

---

### **Week 9: Agriculture Program**

Aim: To analyze agricultural data and extract meaningful insights.

Programs:

### 1. Program 1: Loading Agriculture Dataset

- Procedure: Load the agricultural dataset and display its structure.

- Code:

```
df = pd.read_csv('agriculture_data.csv')  
  
print(df.head())
```

- Expected Result: A preview of the agriculture dataset.

### 2. Program 2: Summary Statistics

- Procedure: Generate summary statistics of the dataset.

- Code:

```
print(df.describe())
```

- Expected Result: Summary statistics including count, mean, std, min, and max values.

### 3. Program 3: Crop Yield Visualization

- Procedure: Visualize crop yields based on different factors.

- Code:

```
sns.boxplot(x='crop_type', y='yield', data=df)  
  
plt.title("Crop Yield Distribution by Type")  
  
plt.xlabel("Crop Type")  
  
plt.ylabel("Yield")  
  
plt.show()
```

- Expected Result: A box plot showing the distribution of yields for different crop types.

### 4. Program 4: Relationship Between Fertilizer Use and Yield

- Procedure: Analyze the relationship between fertilizer use and crop yield.

- Code:

```
sns.scatterplot(x='fertilizer_use', y='yield', data=df)  
  
plt.title("Fertilizer Use vs Crop Yield")  
  
plt.xlabel("Fertilizer Use")  
  
plt.ylabel("Crop Yield")
```

```
plt.show()
```

- Expected Result: A scatter plot showing the relationship between fertilizer use and crop yield.

## 5. Program 5: Predicting Crop Yield

- Procedure: Build a simple linear regression model to predict crop yield.
- Code:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

X = df[['fertilizer_use', 'rainfall']]

y = df['yield']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LinearRegression()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
```

- Expected Result: A mean squared error indicating the prediction error of the model.

---

## Week 10: Marketing Program

Aim: To analyze marketing data and derive actionable insights.

Programs:

### 1. Program 1: Loading Marketing Dataset

- Procedure: Load the marketing dataset and display its structure.
- Code:

```
df = pd.read_csv('marketing_data.csv')

print(df.head())
```

- Expected Result: A preview of the marketing dataset.

### 2. Program 2: Summary Statistics

- Procedure: Generate summary statistics of the dataset.

- Code:

```
print(df.describe())
```

- Expected Result: Summary statistics including count, mean, std, min, and max values.

### 3. Program 3: Customer Segmentation

- Procedure: Segment customers based on their spending behavior.

- Code:

```
from sklearn.cluster import KMeans
```

```
X = df[['annual_spending', 'frequency']]
```

```
kmeans = KMeans(n_clusters=3)
```

```
df['segment'] = kmeans.fit_predict(X)
```

```
sns.scatterplot(x='annual_spending', y='frequency', hue='segment', data=df, palette='Set2')
```

```
plt.title("Customer Segmentation")
```

```
plt.xlabel("Annual Spending")
```

```
plt.ylabel("Frequency")
```

```
plt.show()
```

- Expected Result: A scatter plot showing different customer segments based on spending behavior.

### 4. Program 4: Analyzing Marketing Campaign Success

- Procedure: Analyze the success rate of a marketing campaign.

- Code:

```
campaign_success_rate = df['campaign_success'].value_counts(normalize=True) * 100
```

```
campaign_success_rate.plot(kind='bar', color='green')
```

```
plt.title("Marketing Campaign Success Rate")
```

```
plt.xlabel("Success")
```

```
plt.ylabel("Percentage")
```

```
plt.show()
```

- Expected Result: A bar chart displaying the percentage of successful versus unsuccessful marketing campaigns.

## 5. Program 5: Predicting Customer Churn

- Procedure: Build a logistic regression model to predict customer churn.

- Code:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

X = df.drop('churn', axis=1) # Assuming 'churn' is the target column
y = df['churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LogisticRegression(max_iter=1000)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Accuracy of Churn Prediction Model:", accuracy_score(y_test, y_pred))
```

- Expected Result: An accuracy score indicating the model's performance in predicting customer churn.